

STAT636 - Homework 3

Daniel Osorio - dcosoriorh@tamu.edu

Department of Veterinary Integrative Biosciences

Texas A&M University

1. Consider the Auto data. These represent a two-factor experiment on cars. The two factors are (i) the number of cylinders (4 and 6 were considered) and (ii) origin (three origins were considered). We have 4 cars under each of the $2 \times 3 = 6$ factor combinations. For each car, we have measurements of three weight variables: X_1 = displacement, X_2 = horsepower, and X_3 = acceleration. So, in terms of a two-way MANOVA model, $g = 2$, $b = 3$, and $n = 4$.

```
> auto<- read.csv("Auto_hw.csv")
> auto$cylinders <- as.factor(auto$cylinders)
> auto$origin <- as.factor(auto$origin)
```

- (a) Test for a location effect, a variety effect, and a location-variety interaction at $\alpha = 0.05$. Do this using the manova function in R. Overall, what do you conclude about these data? *As all the p-values are lower than the given α , I can conclude that there is enough evidence against the H_0 , suggesting that there is a real difference in the means associated to the analyzed factors*

```
> manovaResult <- manova(cbind(displacement, horsepower, acceleration) ~
+                          origin + cylinders + origin*cylinders, auto)
> manovaResult
```

Call:

```
manova(cbind(displacement, horsepower, acceleration) ~ origin +
cylinders + origin * cylinders, auto)
```

Terms:

	origin	cylinders	origin:cylinders	Residuals
resp 1	9843.25	34808.17	5766.58	4692.00
resp 2	787.00	4082.04	1057.33	6857.25
resp 3	23.24	5.04	5.74	105.91
Deg. of Freedom	2	1	2	18

Residual standard errors: 16.14517 19.51815 2.425673

Estimated effects may be unbalanced

```
> summary(manovaResult, test = "Wilks")
```

	Df	Wilks	approx F	num Df	den Df	Pr(>F)
origin	2	0.20304	6.503	6	32	0.0001468 ***
cylinders	1	0.11520	40.961	3	16	9.813e-08 ***
origin:cylinders	2	0.25986	5.129	6	32	0.0008635 ***
Residuals	18					

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

- (b) Construct the two-way MANOVA table by computing SSP_{FAC1} , SSP_{FAC2} , SSP_{INT} , SSP_{RES} , and SSP_{COR} . Provide R code that matches the Wilks' statistics computed by manova. Note that

your p-values (computed according to the notes) will not match those of manova, because the distributional results we have learned for two-way MANOVA are large-sample approximations. That said, how do your p-values compare to those of manova? *All the p-values allow performing the same inference about the effects of the factors analyzed. Those computed manually tend to be smaller than the reported by the manova function*

```
> n <- 4
> p <- 3
> g <- 2
> b <- 3
> xBar <- apply(auto[,2:4],2,mean)
> SSPfac1 <- SSPfac2 <- SSPint <- SSPres <- SSPcor <- 0
> for (l in unique(auto$cylinders)){
+   xBar_l <- apply(auto[auto[,1] == l,2:4],2,mean)
+   SSPfac1 <- SSPfac1 + ((b * n * (xBar_l-xBar)^2))
+   for(k in unique(auto$origin)){
+     xBar_k <- apply(auto[auto[,5] == k,2:4],2,mean)
+     xBar_lk <- apply(auto[auto[,1] == l & auto[,5] == k,2:4],2,mean)
+     SSPint <- SSPint + (n * ((xBar_lk - xBar_l - xBar_k + xBar)^2))
+     for(r in seq_len(n)){
+       x <- auto[auto[,1] == l & auto[,5] == k & seq_len(n) == r,2:4]
+       SSPres <- SSPres + ((x - xBar_lk)^2)
+       SSPcor <- SSPcor + ((x - xBar)^2)
+     }
+   }
+ }
> for(k in unique(auto$origin)){
+   xBar_k <- apply(auto[auto[,5] == k,2:4],2,mean)
+   SSPfac2 <- SSPfac2 + (g * n * ((xBar_k - xBar)^2))
+ }
> manovaTable <- cbind(SSPfac1,SSPfac2, SSPint,
+                       t(SSPres), t(SSPcor))
> colnames(manovaTable) <- c("SSP_fac1", "SSP_fac2", "SSP_int",
+                             "SSP_res", "SSP_cor")
> manovaTable
```

	SSP_fac1	SSP_fac2	SSP_int	SSP_res	SSP_cor
displacement	34808.166667	9843.25000	5766.583333	4692.00	55110.0000
horsepower	4082.041667	787.00000	1057.333333	6857.25	12783.6250
acceleration	5.041667	23.24333	5.743333	105.91	139.9383

```
> xBar <- as.numeric(apply(auto[,2:4],2,mean))
> SSPfac1 <- SSPfac2 <- SSPint <- SSPres <- SSPcor <- 0
> for (l in unique(auto$cylinders)){
+   xBar_l <- as.numeric(apply(auto[auto[,1] == l,2:4],2,mean))
+   SSPfac1 <- SSPfac1 + ((b * n * (xBar_l-xBar) %*% t(xBar_l - xBar)))
+   for(k in unique(auto$origin)){
+     xBar_k <- as.numeric(apply(auto[auto[,5] == k,2:4],2,mean))
+     xBar_lk <- as.numeric(apply(auto[auto[,1] == l &
+                                   auto[,5] == k,2:4],2,mean))
+     SSPint <- SSPint + (n * ((xBar_lk - xBar_l - xBar_k + xBar) %*%
+                               t(xBar_lk - xBar_l - xBar_k + xBar)))
+     for(r in seq_len(n)){
```

```

+       x <- as.numeric(auto[auto[,1] == 1 & auto[,5] == k &
+                           seq_len(n) == r,2:4])
+       SSPres <- SSPres + ((x - xBar_1k) %*% t(x - xBar_1k))
+       SSPcor <- SSPcor + ((x - xBar) %*% t(x - xBar))
+     }
+   }
+ }
> for(k in unique(auto$origin)){
+   xBar_k <- apply(auto[auto[,5] == k,2:4],2,mean)
+   SSPfac2 <- SSPfac2 + (g * n * ((xBar_k - xBar) %*% t(xBar_k - xBar)))
+ }
> # Cylinders
> Lambda <- det(SSPres) / det(SSPfac1 + SSPres)
> Lambda
[1] 0.1152037
> 1 - pf((((g * b * (n - 1) - p + 1) / 2) / ((abs((g - 1) - p) + 1) / 2)) *
+ (1 - Lambda) / Lambda, abs((g - 1) - p) + 1, g * b * (n - 1) - p + 1)
[1] 9.813233e-08
> # Origin
> Lambda <- det(SSPres) / det(SSPfac2 + SSPres)
> Lambda
[1] 0.2030373
> 1 - pf((((g * b * (n - 1) - p + 1) / 2) / ((abs((b - 1) - p) + 1) / 2)) *
+ (1 - Lambda) / Lambda, abs((b - 1) - p) + 1, g * b * (n - 1) - p + 1)
[1] 2.888066e-06
> # Interaction
> Lambda <- det(SSPres) / det(SSPint + SSPres)
> Lambda
[1] 0.2598601
> 1 - pf((((g * b * (n - 1) - p + 1) / 2) /
+ ((abs((g - 1) * (b - 1) - p) + 1) / 2)) *
+ (1 - Lambda) / Lambda, abs((g - 1) * (b - 1) - p) +
+ 1, g * b * (n - 1) - p + 1)
[1] 2.079298e-05

```

2. Conduct a simulation study to investigate the coverage probabilities of different confidence interval types with multivariate data. Let the sample size be $n = 30$, the number of variables $p = 5$, the number of simulations $B = 10000$, and

$$\mu' = (0, 0, 0, 0, 0)$$

and

$$\Sigma = \begin{pmatrix} 1.0 & 0.6 & 0.6 & 0.6 & 0.6 \\ 0.6 & 1.0 & 0.6 & 0.6 & 0.6 \\ 0.6 & 0.6 & 1.0 & 0.6 & 0.6 \\ 0.6 & 0.6 & 0.6 & 1.0 & 0.6 \\ 0.6 & 0.6 & 0.6 & 0.6 & 1.0 \end{pmatrix}$$

For each of B times, simulate a dataset of size n from the $N_p(\mu, \Sigma)$ distribution, compute 95% confidence intervals of types one-at-a-time, T^2 simultaneous, and Bonferroni simultaneous, and record

whether each interval contains its corresponding population mean component value. Report a 3×5 matrix of estimated coverage probabilities. The rows of your matrix should correspond to the 3 different interval types, and the columns should correspond to the p mean components; be sure to clearly indicate which row goes with which interval type. Comment on the performance of the different interval types. *The confidence intervals using the quantiles by resampling tend to be more relaxed in detecting differences with respect to the mean, meanwhile, the Bonferroni ones seem to be more accurate. T2 based confidence intervals are between the other two type of intervals in terms of accuracy.*

```
> simulationFunction <- function(n, p, rho, alpha, B = 1000) {
+   mu <- rep(0, p)
+   Sigma <- matrix(rho, nrow = p, ncol = p); diag(Sigma) <- 1
+   F_crit <- (n - 1) * p * qf(1 - alpha, p, n - p) / (n - p)
+   bon_crit <- qt(1 - alpha / (2 * p), n - 1)
+   cov_Q <- cov_T2 <- cov_bon <- matrix(NA, nrow = B, ncol = p)
+   for(i in seq_len(B)) {
+     X <- MASS::mvrnorm(n, mu, Sigma)
+     x_bar <- colMeans(X)
+     S <- var(X)
+     for(k in 1:p) {
+       ci_Q <- quantile(X[,k], c((alpha/2), (1-(alpha/2))))
+       ci_T2 <- x_bar[k] + c(-1, 1) * sqrt(F_crit * S[k, k] / n)
+       ci_bon <- x_bar[k] + c(-1, 1) * bon_crit * sqrt(S[k, k] / n)
+       cov_Q[i, k] <- (ci_Q[1] <= mu[k]) & (mu[k] <= ci_Q[2])
+       cov_T2[i, k] <- (ci_T2[1] <= mu[k]) & (mu[k] <= ci_T2[2])
+       cov_bon[i, k] <- (ci_bon[1] <= mu[k]) & (mu[k] <= ci_bon[2])
+     }
+   }
+   out <- matrix(data = NA, nrow = 3, ncol = p)
+   rownames(out) <- c("95%:", "T2:", "Bonferroni:")
+   out[1,] <- apply(cov_Q, 2, mean)
+   out[2,] <- apply(cov_T2, 2, mean)
+   out[3,] <- apply(cov_bon, 2, mean)
+   return(out)
+ }
> simulationFunction(n = 30, p = 5, alpha = 0.05, rho = 0.6, B = 10000)
```

	[,1]	[,2]	[,3]	[,4]	[,5]
95%:	1.0000	1.0000	1.0000	1.0000	1.0000
T2:	0.9997	0.9996	0.9998	0.9999	0.9996
Bonferroni:	0.9909	0.9912	0.9915	0.9909	0.9908