

STAT636 - Homework 6

Daniel Osorio - dcosoriorh@tamu.edu

Department of Veterinary Integrative Biosciences

Texas A&M University

1. Consider the `Life_Expectancy` data. Let the response Y be `Life.expectancy`, and consider the numeric variables `Income.composition.of.resources`, and `Schooling` as the predictor variables. For the purpose of predicting the value of Y , we will consider non-linear models.

```
> lifeExpectancy <- read.csv("Life_Expectancy.csv")
> lifeExpectancy <- lifeExpectancy[,c("Life.expectancy", "Schooling")]
> lifeExpectancy <- lifeExpectancy[complete.cases(lifeExpectancy),]
```

(a) Polynomial regression.

- i. First predict `Life.expectancy` using a fourth-degree polynomial in `Schooling`. From the p-value of coefficients, what degree of polynomial do you think is necessary? *From the p-values, less than 4.*

```
> lifeExpectancy[,2:5] <- sapply(1:4, function(x){lifeExpectancy[,2]^x})
> colnames(lifeExpectancy) <- c("Y", "X1", "X2", "X3", "X4")
> p4 <- lm(Y~., data = lifeExpectancy[,1:5])
> summary(p4)
```

Call:

```
lm(formula = Y ~ ., data = lifeExpectancy[, 1:5])
```

Residuals:

	Min	1Q	Median	3Q	Max
	-16.6941	-2.3720	0.7387	3.2998	14.2940

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	61.2162412	9.8500202	6.215	3.92e-09 ***
X1	-3.0767145	4.4020089	-0.699	0.486
X2	0.3793709	0.6931948	0.547	0.585
X3	-0.0004507	0.0449784	-0.010	0.992
X4	-0.0004257	0.0010250	-0.415	0.678

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 5.339 on 168 degrees of freedom

Multiple R-squared: 0.6578, Adjusted R-squared: 0.6496

F-statistic: 80.73 on 4 and 168 DF, p-value: < 2.2e-16

- ii. Use a third-degree polynomial in `Schooling` to fit the data again. Then create a grid of values for `Schooling`

```
> school_grid <- seq(from = 1.53, to = 20.04)
```

and make predictions for each. Finally plot the data, and add the fit from the degree-3 polynomial, with 2 times the standard error shown as error bars.

```
> lifeExpectancy <- lifeExpectancy[,1:4]
> p3 <- lm(Y~., data = lifeExpectancy)
> summary(p3)

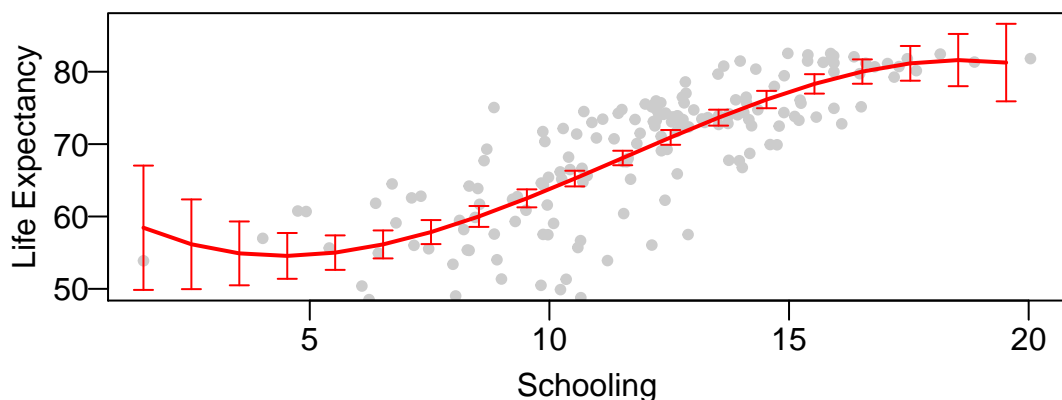
Call:
lm(formula = Y ~ ., data = lifeExpectancy)

Residuals:
    Min       1Q   Median       3Q      Max
-16.8110  -2.2120   0.7029   3.2088  14.2981

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  64.183066   6.764306   9.488  < 2e-16 ***
X1           -4.711970   1.962702  -2.401  0.017447 *
X2             0.656973   0.182969   3.591  0.000432 ***
X3            -0.018992   0.005419  -3.505  0.000585 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 5.326 on 169 degrees of freedom
Multiple R-squared:  0.6574,    Adjusted R-squared:  0.6513
F-statistic: 108.1 on 3 and 169 DF,  p-value: < 2.2e-16

> school_grid <- as.data.frame(sapply(1:3, function(x){school_grid~x}))
> colnames(school_grid) <- c("X1", "X2", "X3")
> sE <- predict(p3,school_grid, se.fit = TRUE)
> par(mar=c(4,4,1,1), mgp = c(1.7,0.5,0))
> x <- school_grid[,1]
> ubSE <- sE$fit + (2* sE$se.fit)
> lbSE <- sE$fit - (2* sE$se.fit)
> yLim <- c(lbSE, ubSE)
> plot(lifeExpectancy[,2:1],
+      pch = 20, col = "gray80", las = 1, xlab= "Schooling",
+      ylab= "Life Expectancy", ylim= c(min(yLim), max(yLim)))
> points(y = sE$fit, x = x, type = "l" ,col= "red", lwd=2)
> arrows(x,ubSE,x,lbSE, length=0.05, angle=90, code=3, col = "red")
```



(b) Splines.

- i. Fit `Life.expectancy` to `Schooling` using a regression model with cubic splines with specified knots at `Schooling` 5, 10 and 15. Make predictions at the grid you created above, and plot the data and add your cubic splines with 2 times the standard error shown as error bars.

```
> library(splines)
> lifeExpectancy <- lifeExpectancy[,1:2]
> b1 <- lm(Y ~ bs(X1, knots = c(5,10,15)), data = lifeExpectancy)
> summary(b1)
```

Call:

```
lm(formula = Y ~ bs(X1, knots = c(5, 10, 15)), data = lifeExpectancy)
```

Residuals:

Min	1Q	Median	3Q	Max
-16.6422	-2.1313	0.8818	3.1406	14.9636

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	53.809	5.317	10.120	< 2e-16 ***
bs(X1, knots = c(5, 10, 15))1	12.293	10.042	1.224	0.222623
bs(X1, knots = c(5, 10, 15))2	-2.830	6.644	-0.426	0.670643
bs(X1, knots = c(5, 10, 15))3	8.877	5.866	1.513	0.132146
bs(X1, knots = c(5, 10, 15))4	26.251	5.916	4.438	1.65e-05 ***
bs(X1, knots = c(5, 10, 15))5	27.046	6.368	4.247	3.60e-05 ***
bs(X1, knots = c(5, 10, 15))6	28.522	7.219	3.951	0.000115 ***


Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 5.318 on 166 degrees of freedom

Multiple R-squared: 0.6646, Adjusted R-squared: 0.6525

F-statistic: 54.82 on 6 and 166 DF, p-value: < 2.2e-16

```
> X1 <- school_grid[,1]
> sE <- predict(b1,bs(X1, knots = c(5,10,15)), se.fit = TRUE)
> x <- school_grid[,1]
> ubSE <- sE$fit + (2* sE$se.fit)
> lbSE <- sE$fit - (2* sE$se.fit)
> par(mar=c(4,4,1,1), mgp = c(1.7,0.5,0))
> plot(lifeExpectancy[,2:1], pch = 20, col = "gray80", las = 1,
+      xlab= "Schooling", ylab= "Life Expectancy",
+      ylim= c(min(c(lbSE, ubSE)), max(c(lbSE, ubSE))))
> points(y = sE$fit, x = x, type = "l", col= "red", lwd=2)
> arrows(x,ubSE,x,lbSE, length=0.05, angle=90, code=3, col = "red")
```



- ii. Fit `Life.expectancy` to `Schooling` using a regression model with natural splines with degree of freedom 4. Make predictions at the grid you created above, and plot the data and add your cubic splines with 2 times the standard error shown as error bars. Note that at the boundary of data, natural splines are better than basic cubic splines.

```
> library(splines)
> b2 <- lm(Y ~ ns(X1, df = 4), data = lifeExpectancy)
> summary(b2)

Call:
lm(formula = Y ~ ns(X1, df = 4), data = lifeExpectancy)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-16.5018	-2.2104	0.8687	3.2655	15.2364

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	58.039	3.789	15.316	< 2e-16 ***
ns(X1, df = 4)1	13.319	3.648	3.651	0.000348 ***
ns(X1, df = 4)2	21.800	2.728	7.991	2.04e-13 ***
ns(X1, df = 4)3	19.085	8.430	2.264	0.024847 *
ns(X1, df = 4)4	28.507	3.715	7.673	1.30e-12 ***


Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 5.267 on 168 degrees of freedom

Multiple R-squared: 0.6669, Adjusted R-squared: 0.659

F-statistic: 84.09 on 4 and 168 DF, p-value: < 2.2e-16

```
> sE <- predict(b2, ns(X1, df = 4), se.fit = TRUE)
> ubSE <- sE$fit + (2* sE$se.fit)
> lbSE <- sE$fit - (2* sE$se.fit)
> par(mar=c(4,4,1,1), mgp = c(1.7,0.5,0))
> plot(lifeExpectancy[,2:1],
+      pch = 20, col = "gray80",
+      las = 1, xlab= "Schooling", ylab= "Life Expectancy",
+      ylim= c(min(c(lbSE, ubSE)), max(c(lbSE, ubSE))))
> points(y = sE$fit, x = x, type = "l", col= "red", lwd=2)
> arrows(x,ubSE,x,lbSE, length=0.05, angle=90, code=3, col = "red")
```



- iii. For smoothing splines, select the λ by cross-validation; what's your result? Also plot the data and add your smoothing splines.

```
> lifeExpectancy <- unique(lifeExpectancy)
> b3 <- smooth.spline(x = lifeExpectancy[,2],
+                     y = lifeExpectancy[,1],
+                     cv = TRUE)
> b3
```

Call:

```
smooth.spline(x = lifeExpectancy[, 2], y = lifeExpectancy[, 1],
              cv = TRUE)
```


Smoothing Parameter spar= 1.072785 lambda= 0.004107205 (17 iterations)

Equivalent Degrees of Freedom (Df): 5.57255

Penalized Criterion (RSS): 4382.429

PRESS(1.o.o. CV): 28.2348

```
> par(mar=c(4,4,1,1), mgp = c(1.7,0.5,0))
> plot(lifeExpectancy[,2:1],
+      pch = 20, col = "gray80", las = 1,
+      xlab= "Schooling", ylab= "Life Expectancy")
> points(x = b3$x, y = b3$y, type = "l", col = "red", lwd=2)
```




2. Consider the `wine_quality_red` data. These contain eleven different variables (1 : fixed acidity; 2 : volatile acidity; 3 : citric acid; 4 : residual sugar; 5 : chlorides; 6 : free sulfur dioxide; 7 : total sulfur dioxide; 8 : density; 9 : pH; 10 : sulphates; 11 : alcohol) and one output (quality (score between 0 and 10)). Use `tree` method and `random forest` method to perform classification.

```
> wineData <- read.csv("winequality_red.csv")
```

(a) Use tree method to do classification, and plot the tree graph

```
> treeModel <- tree::tree(quality~., wineData)
> par(mar=c(1,1,1,1))
> plot(treeModel)
> text(treeModel)
```



i. What's the training error rate?

```
> mean((wineData$quality - predict(treeModel))^2)
[1] 0.4254908
```

ii. Randomly sample 1000 rows as training data and use tree method to do classification. What's the accuracy rate for test data? (Do `set.seed(1)` before sampling.)

```
> set.seed(1)
> n <- seq_len(nrow(wineData))
> training <- sample(n, size = 1000)
> testing <- n[!n %in% training]
> training <- wineData[training,]
> testing <- wineData[testing,]
> treeModel <- rpart::rpart(quality~., training)
> prediction <- predict(treeModel, newdata = testing[, -12])
> mean((testing$quality - prediction)^2)
[1] 0.4353023
```

(b) Use a random forest to do classification with the same training data set. At each split, allow 7 variables to be considered. (Do `set.seed(1)` before doing random forest.)

```
> set.seed(1)
> rForest <- randomForest::randomForest(quality~., training, mtry = 7)
```

i. What's the accuracy rate now?

```
> rForest
```

Call:

```
randomForest(formula = quality ~ ., data = training, mtry = 7)
```

```
  Type of random forest: regression
```

```
    Number of trees: 500
```

```
No. of variables tried at each split: 7
```


Mean of squared residuals: 0.3818267

% Var explained: 41.82

- ii. Use `varImpPlot` to plot the importance figure. What's the most important factor? *Alcohol*

```
> par(mar=c(4,4,1,1), mgp = c(1.7,0.5,0))
```

```
> randomForest::varImpPlot(rForest)
```



- iii. Use bootstrap (100 times) to estimate the standard deviation of classification accuracies. (`set.seed(1)` before bootstrap)

```
> set.seed(1)
```

```
> MSE <- sapply(1:100, function(x){
+   selectedR <- sample(seq_len(nrow(testing)), replace = TRUE)
+   prediction <- predict(rForest, testing[selectedR, 1:11])
+   mean((testing[selectedR, 12] - prediction)^2)
+ })
```

```
> sd(MSE)
```

```
[1] 0.02706819
```

- (c) Consider the binary classification problem. Suppose we denote the wine whose quality is larger than 5 as 'good', otherwise we think it is 'bad'. Similarly as before, randomly sample 1000 observations as training set, and use the tree method to do classification again. (`set.seed(1)` before sampling)

```
> wineData[,12] <- factor(ifelse(test = wineData[,12] > 5,
+                               yes = "G",
+                               no = "B"))
```

```
> set.seed(1)
```

```
> training <- sample(n, size = 1000)
```

```
> testing <- n[!n %in% training]
```

```
> training <- wineData[training,]
```

```
> testing <- wineData[testing,]
```

```
> bC <- tree::tree(quality~., training)
```

- i. What's the accuracy rate, sensitivity and specificity?

```
> ConfussionMatrix <- table(Observed=testing[,12],
+   Predicted=predict(bC,testing[, -12], type = "class"))
> ConfussionMatrix
```

```
      Predicted
Observed  B   G
      B 225  35
      G 127 212
```

```

> # Accuracy
> sum(diag(ConfussionMatrix)/sum(ConfussionMatrix))
[1] 0.7295492
> # Sensitivity (True Positive Rate)
> ConfussionMatrix[2,2]/sum(ConfussionMatrix[2,])
[1] 0.6253687
> # Specificity (True Negative Rate)
> ConfussionMatrix[1,1]/sum(ConfussionMatrix[1,])
[1] 0.8653846


```

- ii. Let $\kappa = \text{seq}(\text{from} = 0, \text{to} = 1, \text{by} = 0.01)$. Plot the ROC curve.

```

> bProb <- predict(bC,testing[, -12])[,1]
> K <- sapply(seq(from = 0, to = 1, by = 0.01), function(x){
+   ifelse(test = bProb > x, yes = "B", no = "G")
+ })
> ROC <- t(apply(K, 2, function(hat){
+   ConfussionMatrix <- table(Observed=testing[,12],
+                             Predicted=factor(hat, levels = c("B", "G")))
+   c(ConfussionMatrix[1,2]/sum(ConfussionMatrix[1,]),
+     ConfussionMatrix[2,2]/sum(ConfussionMatrix[2,]))
+ })))
> par(mar=c(4,4,1,1), mgp=c(1.6,0.5,0))
> plot(ROC, type = "l", ylab = "True Positive Rate",
+      xlab = "False Positive Rate", main = "ROC", las = 1)

```



- iii. Use cross validation to decide the number of nodes. What's the best choice? (`set.seed(1)` before cross validation)

```

> accNodes <- c(NA, sapply(2:12, function(nNumber){
+   hat <- predict(tree::prune.tree(bC, best = nNumber),
+     testing[, -12], "class")
+   ConfussionMatrix <- table(Observed=testing[,12],
+     Predicted=factor(hat, levels = c("B", "G")))
+   sum(diag(ConfussionMatrix)/sum(ConfussionMatrix))
+ })))
> accNodes
[1] NA 0.7061770 0.7061770 0.7061770 0.7061770 0.7345576 0.7378965
[8] 0.7378965 0.7445743 0.7312187 0.7312187 0.7295492
> which.max(accNodes)

```


[1] 9

- iv. 'Prune' your tree with respect to the number of nodes as 4, 7, 8 and plot the ROC curves in one figure with your original ROC curve in (c) which contains 9 nodes. From the figure, how many nodes would you choose?

```
> par(mar=c(4,4,1,1), mgp = c(1.7,0.5,0))
> plot(ROC, type = "l", ylab = "True Positive Rate",
+      xlab = "False Positive Rate", main = "ROC", las = 1)
> colN <- 1
> for(i in c(4,7,8)){
+   bCp <- tree::prune.tree(bC,best = i)
+   bProb <- predict(bCp,testing[,12])[,1]
+   K <- sapply(seq(from = 0, to = 1, by = 0.01), function(x){
+     ifelse(test = bProb > x,yes = "B", no = "G")
+   })
+   ROC <- t(apply(K,2,function(hat){
+     ConfussionMatrix <- table(Observed=testing[,12],
+                               Predicted=factor(hat, levels = c("B","G")))
+     c(ConfussionMatrix[1,2]/sum(ConfussionMatrix[1,]),
+       ConfussionMatrix[2,2]/sum(ConfussionMatrix[2,]))
+   })))
+   par(mar=c(4,4,1,1), mgp=c(1.6,0.5,0))
+   colN <- colN + 1
+   lines(ROC, lty = colN, col = colN)
+ }
> legend("bottomright",
+       legend = c(4,7,8,9),
+       lty = c(2,3,4,1),
+       col = c(2,3,4,1),
+       bty = "n")
```

